

Sentinel

A Privacy-Preserving Edge-AI Architecture for Proactive Mental Health Anomaly Detection

Anant

B.Tech, Computer Science and Engineering (Data Science)

[Institution name, department, and guide/supervisor details to be added]

Abstract

More than one billion people currently live with a diagnosable mental health condition, and suicide remains among the leading causes of death in adolescents and young adults worldwide. Passive smartphone sensing — termed digital phenotyping — has been proposed as a scalable complement to episodic clinical assessment, since continuously generated behavioural signals (communication patterns, circadian rhythm, mobility, application usage) can reveal symptom trajectories between clinical visits. However, existing digital-phenotyping research infrastructure and consumer wellness features rely on cloud-mediated processing or centralized data aggregation, creating a structural conflict between clinical utility and user privacy that has limited real-world, at-scale deployment. This paper presents Sentinel, an architecture for proactive mental health anomaly detection designed as a conceptual extension to the Android Digital Wellbeing framework, under a strict zero-knowledge constraint: no raw behavioural data — text, media, location, or usage logs — ever leaves the device. Sentinel couples a two-phase on-device TensorFlow Lite pipeline (lightweight NLP sentiment/risk scoring via TF-IDF and MobileBERT variants, Gemma 4 followed by an ensemble anomaly-detection layer wrapped in conformal calibration) with data-at-rest encryption via SQLCipher and hardware-backed key storage in the Android Trusted Execution Environment. Generation of active, conversational content is grounded in clinical instruments via a retrieval-augmented mechanism and verified by a dedicated safeguarding layer before reaching the user. Critically, Sentinel replaces automated third-party notification with a graduated, human-anchored intervention protocol, escalating to existing human-staffed crisis infrastructure by default, with any external contact reserved for a user-preconfigured, explicitly opted-in mechanism. We describe the system architecture, examine the computational feasibility of continuous on-device inference under mobile hardware constraints, propose a staged validation plan extending existing evaluation frameworks to the detection side of the system, and outline an ethical framework addressing false-positive harm, validation without clinical ground truth, cultural and interpretive robustness, and the risk of intervention-protocol misuse.

1. Introduction

Mental health disorders now represent one of the most pressing challenges in global public health. World Health Organization estimates place over one billion people worldwide as living with a mental health condition, with anxiety and depressive disorders the most prevalent across both sexes. Suicide alone accounted for an estimated 727,000 deaths in 2021 and remains the third leading cause of death among 15–29 year-olds globally [2]. Yet clinical detection of deteriorating mental state continues to rely overwhelmingly on episodic, self-reported assessment — structured interviews and standardized questionnaires administered at intervals of weeks or months — a cadence poorly matched to the temporal dynamics of mood and suicidal ideation, which can fluctuate meaningfully within a single day.

Smartphones, carried continuously and instrumented with an increasingly rich sensor suite, offer a candidate solution. The resulting research paradigm — digital phenotyping, defined as the moment-by-moment quantification of individual-level human behavior using data from personal digital devices [1] — has motivated over a decade of work correlating passively sensed behavioral markers with depressive symptom severity. Saeb et al. found that GPS-derived features (variance in locations visited, regularity of 24-hour movement rhythm) and phone-usage frequency correlated significantly with depression severity over a two-week sensing period across 40 adults [3]. The StudentLife study similarly found that accelerometer-derived stationary time and screen-unlock duration distinguished depressed from non-depressed students across a full academic term, in a cohort of 83 participants [4]. Open research platforms such as Beiwe, developed at the Harvard T.H. Chan School of Public Health [1], and mindLAMP, developed at Beth Israel Deaconess Medical Center / Harvard Medical School [5], have operationalized this paradigm at scale, coupling passive sensor collection with clinician-facing dashboards for research and, increasingly, clinical use.

This clinical promise has not translated into consumer-scale deployment, largely because existing digital-phenotyping infrastructure was built for research data collection rather than end-user privacy. Beiwe, mindLAMP, and comparable platforms are explicitly designed to transmit sensor and survey data to a server for aggregation and clinician review — an appropriate model for IRB-governed research, but a poor fit for an OS-level consumer feature where the user, not a research team, must remain the sole party with access to their own behavioral data. The broader machine learning community has converged on a partial answer to this tension: federated and on-device learning architectures, which keep raw data on the client device and transmit only model updates or, in the fully local case, nothing at all [6]. What remains largely unaddressed is combining this privacy posture with the specific, high-stakes task of detecting acute mental health risk, where the cost of a missed detection and the cost of a false positive are both severe and asymmetric.

Intervention design compounds this tension. Prior systems for detecting self-harm and suicide risk in user-generated text — including recent work applying transformer-based classifiers to private social-media conversations, built specifically to distinguish casual expression from genuine risk [7] — have generally

been built for platform-side content moderation, where a detected signal triggers a report to human moderators. This design assumes a trust relationship between the detecting system and any party it notifies; that assumption does not hold for an OS-level, self-monitoring feature, where a poorly designed escalation path can itself become a vector for harm. This is not a hypothetical concern: a substantial body of HCI security research documents how ostensibly protective device features — shared location, shared passwords, “for your safety” monitoring — are routinely repurposed by abusive partners for coercive control [8]. A pre-authorized emergency contact is exactly this shape of feature. To our knowledge, no existing system combines (i) fully local, multimodal behavioral anomaly detection integrated at the operating-system level, (ii) hardware-backed, zero-knowledge data-at-rest encryption, and (iii) a graduated, human-in-the-loop intervention protocol that defaults to user-directed crisis resourcing rather than automated third-party notification. This is the gap Sentinel is designed to address.

This paper makes three contributions. First, we present the Sentinel system architecture: a two-phase, on-device TensorFlow Lite pipeline for behavioral feature extraction and unsupervised anomaly detection, integrated with Android’s native usage, activity-recognition, and Trusted Execution Environment APIs, together with a retrieval- and instrument-grounded generation layer for active conversational check-ins. Second, we examine the computational feasibility of this architecture under real mobile hardware constraints, including the specific limitations of on-device vision models for the more exploratory task of detecting self-harm-related themes in user-generated media. Third, we propose an ethical framework and staged validation plan for the intervention protocol, addressing false-positive harm mitigation, user autonomy, interpretive and cultural robustness, and the risk of protocol misuse in coercive-control contexts. Section 2 reviews related work across digital phenotyping, conversational agents, generation safety, and prediction under uncertainty; Section 3 details the proposed methodology; Section 4 examines implementation feasibility; Section 5 develops the ethical framework; and Section 6 concludes with directions for future validation work.

2. Literature Review

Sentinel sits at the intersection of six research threads: digital phenotyping and passive sensing, conversational agents in mental health, the emerging empathy-paradox literature on what these systems actually understand, generation-safety techniques including retrieval-augmented generation, the statistics of predicting rare events, and the regulatory and ethical scaffolding now forming around clinical AI. This section reviews each in turn and closes by positioning Sentinel against the gap they collectively leave open.

2.1 Digital Phenotyping: From Passive Sensing to Behavioral Prediction

The term digital phenotyping was introduced by Torous, Kiang, Lorme, and Onnela [1], who in the same paper released Beiwe, an open research platform for passive smartphone sensing. Saeb et al. [3] gave this premise its first strong empirical footing, showing that GPS-derived and phone-usage features correlated significantly with depression severity. The StudentLife project [4] extended this to a full academic term, finding that accelerometer-derived stationary time and screen-unlock duration distinguished depressed from non-depressed participants. Both studies converge on the finding that motivates Sentinel’s own behavioral feature set: sleep, mobility, and app-engagement signals are the strongest and most consistently replicated correlates in this literature — stronger than anything derived from media content.

Beiwe’s later counterpart, mindLAMP [5], operationalized this at greater scale by pairing passive sensing with brief active check-ins (Ecological Momentary Assessment) and clinician-facing dashboards. Both platforms share a structural feature that matters for Sentinel’s design: they transmit sensor and survey data to a server for clinician review, appropriate for IRB-governed research but a poor fit for a consumer-facing, zero-knowledge feature. The broader machine learning field has partially answered this tension through federated and on-device learning, formalized by McMahan et al. [6], which keeps raw data on the client device and transmits only model updates, or nothing at all in the fully local case — the case Sentinel adopts as a hard constraint.

One further signal resolves a tension in Sentinel’s own privacy posture: keystroke dynamics. The BiAffect research program [20, 21] demonstrated, across a bipolar cohort and a large open-science iOS sample, that typing metadata alone — inter-key timing, backspace rate, session duration and variability — correlates with clinician-administered mood severity, without ever needing to process the content of what was typed. This is a stricter privacy standard than Sentinel’s text-sentiment layer and is incorporated into the feature set for that reason.

2.2 Conversational Agents in Mental Health Care

The idea that a machine could deliver therapeutic conversation predates modern NLP by half a century: Weizenbaum’s ELIZA (1966) simulated a Rogerian therapist through simple pattern matching, and users nonetheless disclosed intimate material to it — an early instance of what human-computer interaction

research later formalized as the Computers Are Social Actors paradigm. Modern systems have made this dynamic clinically actionable. Woebot, evaluated in a randomized controlled trial by Fitzpatrick, Darcy, and Vierhile [17], delivered CBT to young adults with depression and anxiety symptoms with measurable symptom reduction. Wysa, evaluated by Inkster et al. and summarized in Dehbozorgi et al.'s 2025 systematic review [49], showed that among self-reporting users, higher engagement predicted significantly larger symptom improvement, with 67.7% of users rating the app helpful — an early demonstration that engagement depth, not mere access, drives outcome, the logic behind Sentinel's graduated tiers rather than a single flat intervention.

Not every result in this literature is encouraging in the way marketing materials suggest, and the honest ones are more useful to Sentinel than the flattering ones. Tate et al., also reviewed in Dehbozorgi et al. [49], built a machine learning model to predict adolescent mental health problems from a large Swedish twin cohort and achieved their best result — AUC 0.739 with a random forest — concluding explicitly that the model, despite being the strongest configuration tested, was not accurate enough for clinical use. This is a rare and valuable admission in a literature that Dehbozorgi et al.'s own review characterizes as frequently suffering from inadequate transparency in reporting model features and an overly optimistic view of AI performance — a field-level critique this paper takes as a standard to write against.

Kellogg and Sadeh-Sharvit [53] offer a complementary, practice-oriented lens: rather than asking whether AI can replace clinicians, they organize mental-health AI into three categories clinicians actually interact with — automation, engagement, and clinical decision-support technologies — and document the real friction of integrating each into daily practice. This framework maps cleanly onto Sentinel's own architecture and is used in Section 3 to justify why the human-counsellor handoff needs its own onboarding, not only user-facing design.

2.3 The Empathy Paradox and the Boundary Between Expression and Interpretation

A 2026 study by Kehinde et al. [10] extracted 1,000 AI-directed and 1,000 human-therapist-directed mental health conversations and used a Random Forest classifier to distinguish them by emotional content across five domains. The classifier achieved 48.3% accuracy — chance level — with the modest features that did distinguish the groups being structural (word count) rather than emotional. Users, in other words, express distress to AI systems with the same linguistic and emotional texture used with human clinicians. This finding is easy to over-read: it establishes that human expression toward AI is genuine, not that AI interpretation of that expression is reliable. These are different claims, and the literature on the second is considerably more cautionary.

The Thera-Turing Test [12] makes this distinction structurally explicit: its evaluation framework separates “therapeutic relationship skills” (is the agent empathic, rapport-building) from “treatment fidelity” and “safety” (is the agent's underlying clinical judgment correct), and none of its four evaluation methods specifically stress-tests interpretation under deliberately ambiguous or adversarial input. That gap is not

hypothetical. EmoAgent [11], comprising EmoEval and EmoGuard, simulated vulnerable users interacting with unguarded conversational agents and found psychological deterioration in more than a third of interactions, with delusion-exacerbation rates exceeding 90% for certain agent personas — a documented case of a system sounding appropriately warm while actively worsening the underlying clinical picture. A parallel literature on AI-associated sycophancy makes the same point with more clinical weight: a documented industry rollback of a large conversational model update specifically for validating user doubts and reinforcing negative emotion, and the “Technological folie à deux” analysis of chatbot–mental-illness feedback loops [31], both document sycophantic validation as an active harm mechanism, not a benign stylistic quirk.

A second, distinct axis of interpretive risk is cultural and linguistic. Kleinman’s foundational work on idioms of distress [32], extended by Kohrt and Hruschka’s demonstration that expressions such as the Nepali man dukheko do not map cleanly onto Western depression constructs [33], establishes that culturally coded expressions of distress can be systematically misread by systems calibrated to Western/English norms. This concern is current and active: PsychiatryBench [34] was built explicitly to test whether large language models apply ethnocentric assumptions when interpreting culturally mediated clinical language, and a companion effort on cross-lingual mental health ontologies for Indian languages [35] addresses the same problem in the linguistic context most relevant to Sentinel’s integration with Tele MANAS [14], which itself operates across 20 languages. Section 5.4’s interpretive-robustness battery is a direct methodological response to this literature.

2.4 Grounding Generation: Retrieval-Augmented Generation and Process Knowledge

Retrieval-augmented generation, introduced by Lewis et al. [22] as a technique for grounding language model output in externally retrieved knowledge rather than parametric memory alone, has become the standard first-line mitigation for hallucination in high-stakes generative applications, with domain-specific extensions such as MEGA-RAG [23] adapting the technique for public-health contexts. Amugongo et al.’s 2025 systematic review of RAG in healthcare [54] is the most comprehensive account of where this technique stands: RAG measurably reduces — but does not eliminate — hallucination, with the degree of reduction varying by underlying model. More consequential for Sentinel’s design is what the review found missing: of the healthcare RAG studies that addressed safety at all, only three of thirty-seven evaluated their systems against intentional or unintentional harm, and the majority did not assess privacy or bias exposure, despite the retrieval corpus itself being a documented vector for exposing sensitive information. RAG is necessary but not sufficient, and the review’s own recommended mitigations — robust privacy measures, transparency, bias mitigation, human oversight — describe almost exactly the layered architecture Sentinel already proposes independently of this literature.

Roy et al.’s ProKnow system [52] operationalizes grounded generation for mental health specifically, and is the closest existing precedent to Sentinel’s Tier 1 design. Rather than retrieving from a generic wellness corpus, ProKnow constrains question generation using semantic lexicons built directly from the PHQ-9

and GAD-7 instrument text, refined with clinician input, so that generated follow-up questions remain traceable to established diagnostic criteria. The reported results are the strongest in this literature: constrained generation was 89% safer than unconstrained baseline language models in the depression and anxiety domain, produced a 96% reduction in a knowledge-capture-violation metric, and yielded an 82% average improvement across safety, explainability, and process-adherence combined. The paper also documents, as motivating context, a real case in which unguarded use of a large language model for preventive care produced the response “I think you should” to a user’s message “Should I kill myself?” — a sourced, concrete instance of exactly the failure this paper’s safety architecture, from EmoGuard to the interpretive-robustness battery, exists to prevent. Sentinel’s RAG layer adopts ProKnow’s mechanism directly.

2.5 Prediction Under Uncertainty: The Base-Rate Problem and Calibration

The most consequential finding for how Sentinel frames its own claims comes from outside the AI literature entirely. Franklin et al.’s meta-analysis of 365 studies spanning fifty years of suicide-risk research [18] found that risk-factor-based prediction performed only marginally better than chance across every outcome measure examined, with no improvement across five decades. A related synthesis [19] sharpens why: machine learning suicide-prediction models can achieve strong overall classification accuracy while their accuracy at predicting an actual future suicide event approaches zero — a direct consequence of suicide’s low base rate, not a deficiency correctable with more data or a larger model. Tate et al.’s AUC 0.739 ceiling, discussed in Section 2.2, is a recent, concrete instance of this constraint operating in practice.

The correct response to a hard statistical limit is not to ignore it but to build around it. Using Fitbit-derived heart rate, step count, and sleep data from 13,835 participants in the NIH All of Us dataset, a 2025 study applied conformal prediction to psychiatric condition classification [24]: the raw model achieved a modest 64% mean accuracy, but wrapped in conformal calibration, it produced prediction sets with a 90% coverage guarantee by explicit statistical construction — an honest range rather than an overconfident point estimate. This is the direct precedent for Sentinel’s own conformal-calibration layer (Section 3), applied to a data modality close enough to Sentinel’s own that the pattern transfers.

2.6 Regulatory and Ethical Frameworks for AI in Healthcare

Two recent reviews map the ethical terrain comprehensively enough to serve as a checklist against Sentinel’s own Section 5. Corfmat, Martineau, and Régis [46] organize AI-in-healthcare ethics and law into six categories — privacy, individual autonomy, bias, responsibility and liability, evaluation and oversight, and professional impact — and their responsibility/liability analysis exposes a genuine open question for Sentinel: standard fault-based liability requires establishing duty, breach, and causal harm, a framework built around a human practitioner that does not map cleanly onto a Tier 1 system acting autonomously by design. Saeidnia et al.’s systematic review of 51 articles [47] converges on an overlapping but

independently derived set of considerations, and surfaces the single most directly relevant precedent found in this literature: Mörch, Gupta, and Mishara’s Canada Protocol [48], an ethics checklist built specifically for AI in suicide prevention and mental health.

On evaluation methodology, two further sources give Sentinel’s validation plan real machinery. Rosenthal, Beecy, and Sabuncu [44] contrast the traditional linear model of AI deployment — train, freeze, deploy, monitor — with a “dynamic deployment” model built around continuous feedback and staged validation, and report that more than 40% of a sample of 521 FDA-approved medical AI devices lacked any clinical validation data at all, a concrete warning against exactly the overclaiming this paper is structured to avoid; they also point to the CONSORT-AI and SPIRIT-AI extensions [50, 51] as the actual formal reporting standards for AI clinical trials. Livingston et al. [45] operationalize clinician-in-the-loop evaluation directly: a five-dimension scoring framework applied by independent expert reviewers with mode-based or Delphi consensus resolution, and AHRQ-based severity grading for anything flagged harmful. Their own reported inter-reviewer agreement rate — 60.6% on the first pass, even among trained experts — is itself an argument for calibrated uncertainty over false confidence, echoing Section 2.5 from an entirely different empirical angle.

Finally, two sources ground the paper’s escalation design in real infrastructure rather than an invented mechanism. Freed et al.’s study of intimate partner surveillance [8] documents, empirically, how ostensibly protective device features are routinely repurposed for coercive control — directly informing the mitigation built into Section 5.6’s pre-authorized-contact design. Tele MANAS [14], India’s national tele-mental-health program, already operates the exact function Section 5.2’s escalation protocol depends on: trained human counsellors, not an algorithm, making the acute-risk judgment call.

2.7 Synthesis: The Gap Sentinel Addresses

Three threads run through this literature and converge on the same unaddressed space. First, digital phenotyping research has repeatedly shown that passive behavioral signal correlates meaningfully with mental state, but the platforms that established this — Beiwe, mindLAMP — were built for research-team access, not zero-knowledge consumer deployment. Second, conversational-agent research has demonstrated real therapeutic value alongside an increasingly urgent caution that apparent empathy and sycophantic validation are not the same as correct interpretation, and that this gap is measurable and clinically consequential rather than speculative. Third, the prediction and evaluation literature — from Franklin et al.’s fifty-year null result to Livingston et al.’s expert-disagreement rate — converges on calibrated uncertainty and staged, human-anchored evaluation as the only defensible response to a domain where point-accuracy claims are not statistically supportable. No system reviewed here combines fully local multimodal sensing, hardware-backed zero-knowledge storage, generation grounded in clinical instruments via a ProKnow-style mechanism, conformal-calibrated detection, and a graduated escalation protocol anchored in existing human-staffed infrastructure rather than automated dispatch. That combination is Sentinel’s contribution.

3. Proposed Methodology

3.1 System Overview

Sentinel operates as a five-stage local pipeline: (1) passive behavioral feature acquisition across four modalities, (2) on-device feature extraction via TensorFlow Lite, (3) hardware-encrypted local persistence, (4) ensemble anomaly detection wrapped in conformal calibration, and (5) a graduated, human-anchored intervention protocol. No stage transmits raw data off-device. The single exception in the entire architecture is Tier 3 of the intervention protocol, and even there, only with explicit prior user authorization and only a pre-formatted alert — never raw behavioral data.

3.2 Data Acquisition and Behavioral Feature Streams

Four passive modalities feed the pipeline. Text and typing sentiment is scored locally via MobileBERT, chosen for its bottleneck architecture purpose-built for resource-limited mobile devices [15]. Keystroke dynamics — inter-key timing, backspace rate, session duration and variability — follow the BiAffect methodology [20, 21], capturing metadata about how a person types rather than what they type, a stricter privacy posture than content-based sentiment scoring since it never requires reading message text at all. Circadian and sleep signal is captured through Android’s Sleep API [25] (`SleepSegmentEvent` and `SleepClassifyEvent`), which combines accelerometer and ambient-light data through an on-device classifier under the runtime `ACTIVITY_RECOGNITION` permission. App-usage signal comes from `UsageStatsManager` [27], using `queryEvents()` for granular foreground/background and screen-interactive transitions rather than only aggregated daily totals, gated behind the special `PACKAGE_USAGE_STATS` permission.

A fifth, exploratory modality — vision-based scanning of newly captured media for self-harm-related themes — is deliberately scoped out of the core architecture. The published evidence base for text-based self-harm detection is substantial and growing [7]; the evidence base for lightweight, on-device image classification of comparable nuance is not, and this asymmetry is treated as a design constraint rather than a gap to paper over.

3.3 On-Device Model Architecture

The inference pipeline runs in two phases with fundamentally different resource profiles. Phase 1 is always-on: MobileBERT performs continuous, lightweight sentiment classification on the text stream at negligible battery and thermal cost. Phase 2 is detection: rather than a single Isolation Forest score, an ensemble of Isolation Forest and a rolling statistical baseline (z-score against the personal 7–14 day baseline) runs in parallel, with disagreement between the two treated as a built-in signal of low confidence rather than averaged away. The ensemble output is wrapped in conformal prediction, following the

demonstrated pattern from wearable-derived psychiatric classification [24], so that Sentinel reports a calibrated risk range with a statistical coverage guarantee rather than a single, falsely precise number.

Phase 3 activates only during scheduled or triggered Ecological Momentary Assessment check-ins: a quantized Gemma 4 E2B model (approximately 2 billion active parameters, 4-bit quantized, deployed via ML Kit’s on-device GenAI API) [26], never running continuously. Its output is constrained by two independent safety layers rather than trusted directly. First, retrieval-augmented generation grounds every generated prompt in semantic lexicons built from the PHQ-9 and GAD-7 instrument text — the same clinical instruments that already anchor the EMA design — refined with clinician input, following the mechanism demonstrated by ProKnow to produce measurably safer, more diagnostically traceable generation than unconstrained language models [52]. Second, an EmoGuard-style safeguard layer — an emotion-tracking module, a cognitive-distortion checker, a conversational-boundary module, and a synthesis layer that can override generation — verifies output before it reaches the user [11].

3.4 Secure On-Device Storage

Behavioral data at rest is stored in a local Room database encrypted via SQLCipher (256-bit AES), with cryptographic keys held in the Android Keystore system rather than application memory. Key generation explicitly distinguishes TEE-backed from StrongBox-backed keys [9]: the implementation requests StrongBox via `setIsStrongBoxBacked(true)` where `FEATURE_STRONGBOX_KEYSTORE` is present, falling back to TEE otherwise, since the two provide materially different resistance to physical extraction and are not interchangeable claims. Network permissions are restricted to the single explicit exception described in Section 3.5.

3.5 The Multi-Tiered Intervention Protocol

Detection output routes through three tiers rather than a single threshold, deliberately avoiding both extremes of the design space this paper considered and rejected — a purely autonomous system with no human check, and an automated authority-dispatch mechanism with no realistic implementation path and a documented history of harm.

Tier	Indicators	Sentinel action	Human in the loop
1 — Low risk	Mild stress markers, stable sleep/mobility, minor negative affect	On-device, RAG-grounded CBT-style nudge or resilience exercise	None — bounded to structured, instrument-grounded generation
2 — Medium risk	Clinical-range PHQ-9/GAD-7-style elevation, sustained withdrawal, disrupted sleep	One-tap connection to Tele MANAS Tier 1 (14416) or a user-designated contact	Yes — counsellor or contact makes the judgment call
3 — Critical risk	Signals consistent with suicidal ideation or imminent self-harm	Immediate, pre-dialed surfacing of Tele MANAS / local emergency number; optionally, with prior consent, an	Yes, always — Sentinel never contacts police or dispatches responders unilaterally

		encrypted context message to a pre-designated contact	
--	--	---	--

Tier 2/3 escalations do not surface as raw model output. They surface as a structured summary — which signals fired, their conformal confidence bounds, recent trend — reviewed against two documented human–AI collaboration failure modes, automation bias and alert fatigue, rather than assumed to be self-evidently correct [39, 46]. This mirrors the only close structural precedent identified in this literature, a peer-support system independently proposing the same three-tier logic — autonomous handling for low-stress input, moderated review for medium-stress, direct human intervention for high-stress [40].

Three commitments govern user-facing transparency, following directly from Section 5 rather than existing as separate boilerplate: onboarding discloses in plain language exactly which signals are collected and what each tier does, not merely that “AI monitoring” is active; the user retains standing control to view, adjust, or disable any tier — including the Tier 3 pre-authorized contact — at any time; and the system discloses its non-human status at every active-conversation turn. The pre-authorized contact mechanism is further bounded by a cooling-off period before it can be changed, and a flag if the same contact is repeatedly added and removed, mitigating the documented risk that such features become vectors for coercive control [8].

4. Implementation and Feasibility

4.1 Platform Constraints on Continuous Monitoring

Two Android-specific constraints bound what “always-on” can actually mean in practice. First, UsageStatsManager event-level data is retained by the system for only a few days, so a 7–14 day rolling baseline requires active daily persistence via WorkManager (minimum 15-minute interval) rather than lazy querying. Second, and more consequentially, Android’s standby-bucket system — tightened further under Android 16, which now applies execution quotas even to the “active” bucket — throttles background execution based on how often the monitoring app itself is opened, not on the user’s overall phone engagement. This creates a structural tension worth stating plainly: an app designed to run invisibly, as Sentinel is by design, is a natural candidate for the “rare” standby bucket — capped at roughly ten minutes of background execution per rolling 24 hours — regardless of how engaged the user actually is with their device generally. The mitigation is a persistent low-priority foreground service, which exempts the app from standby throttling at the cost of an always-visible notification — a genuine UX trade-off between reliability and the visible reminder of being monitored, named here rather than assumed away.

4.2 Computational Feasibility of the Model Stack

MobileBERT’s 25.3 million parameters and sub-tens-of-megabyte footprint make continuous Phase 1 inference computationally uncontroversial [15]. The Phase 3 generative component is the harder case: Gemma 4 E2B’s Per-Layer Embedding architecture keeps its active memory footprint to approximately 1.3–2GB in 4-bit quantized form — further reduced below 1GB for text-only inference in its most recent quantization-aware training release — with official Android deployment paths rather than a bespoke integration [26]. This is offered as a specification grounded in current, verified figures, not a measured benchmark; Section 4.3’s Stage 0 explicitly requires on-device latency, memory, and battery measurement on representative mid-range hardware before this claim is treated as demonstrated. The KV-cache scales with context length regardless of model choice, which is why EMA and CBT-nudge interactions are bounded to short context windows by design rather than left open-ended.

Framed against the distinction between linear and dynamic AI deployment [44], Sentinel is deliberately architected as a dynamic system in that sense: the personal baseline is continuously updated from local feedback, and the staged validation plan below builds in recurring re-evaluation rather than a single frozen train-and-deploy cycle. This matters because the same literature reports that more than 40% of a sample of 521 FDA-approved medical AI devices lacked any clinical validation data at all [44] — a documented instance of exactly the shortcut this paper is structured to avoid, and the reason Section 4.3 treats staged validation as load-bearing rather than optional.

4.3 Staged Validation Plan

This paper extends the Thera-Turing Test’s staged-readiness logic [12] — not ready → beta-ready → clinical-trial-ready → full-launch-ready, originally built for conversational-quality assessment — to the detection side of the architecture, since no existing framework covers both.

Stage	Purpose	Method	Gate
0 — Offline / simulation	Validate detection math pre-deployment	Replicate published sensor–mood correlations [3, 4] on public/synthetic data; verify conformal coverage; static/dynamic network analysis for zero data egress; run the interpretive-robustness battery offline; benchmark model latency/memory/battery on mid-range hardware	Coverage holds; zero egress confirmed; battery passed; measured performance within spec
1 — Friendly beta	Real-world false-alarm rate, UX, battery cost	Small opt-in cohort; retention tracked against established smartphone-intervention attrition benchmarks [42]; standby-bucket telemetry	Acceptable false-alarm rate; no nuisance-driven uninstalls; no privacy incidents
2 — Resistant beta + simulated crisis (IRB-reviewed)	Safety under adversarial/ambiguous input	Red-team using expert-validated ethical stress-test vignettes [38]; two independent reviewers score each flagged interaction across defined dimensions, disagreement resolved by mode-of-three or Delphi consensus, harm severity graded by AHRQ classification [45]	Zero missed simulated high-risk cases; zero harmful sycophantic output; correct Tele MANAS routing
3 — Clinical pilot (future work)	Real accuracy/calibration against clinical outcomes	IRB-approved longitudinal study under licensed clinician oversight, mirroring the Thera-Turing Test’s judge criteria [12]	Only this stage could support a clinical-utility claim; Sentinel as presented here does not reach it

Reporting at any stage beyond Stage 0 should follow the CONSORT-AI and SPIRIT-AI extensions [50, 51] — the actual formal standards for AI clinical trial protocols and reports — rather than an invented substitute.

4.4 Anticipated Regulatory Context

Under FDA’s AI/ML-Based Software as a Medical Device Action Plan [43], the large majority of currently authorized AI/ML SaMD enter the market as Class II devices via the 510(k) pathway, with De Novo and PMA — typically Class III — reserved for higher-risk cases. FDA guidance further indicates that AI offering only general health education or lifestyle advice without specific clinical recommendations may fall under enforcement discretion rather than formal SaMD regulation. Under this framework, Sentinel’s Tier 1 plausibly sits outside formal SaMD regulation; Tier 2/3, which involve risk classification and recommended action, would more likely require at minimum 510(k) engagement, with Tier 3 approaching Class III given FDA’s own threshold for software whose failure could cause death or irreversible harm. This is offered strictly as anticipated regulatory context, not a compliance claim — a real deployment would need to engage FDA’s Digital Health Center of Excellence, or India’s CDSCO for domestic deployment, directly rather than rely on this secondary analysis.

5. Ethical Considerations

5.1 The Zero-Knowledge Boundary

Privacy claims (no raw data egress) are treated as a Stage 0 validation target — verified through network-traffic and static analysis — not an assumption embedded in the architecture diagram. A design that intends zero egress and a system that has been shown to achieve it are different claims, and this paper makes only the first until Stage 0 is complete.

5.2 Escalation Policy and the Clinician Review Pathway

Sentinel’s tiers route by risk level, but tiering alone does not specify who reviews an escalation or how. A clinician-in-the-loop review layer is proposed at Tier 2/3: escalations surface as a structured summary — which signals fired, their conformal confidence bounds, recent trend — reviewed against known human-AI collaboration failure modes: confirmation bias, automation bias, and alert fatigue, which recent simulation-based work identifies as the actual determinants of whether clinicians appropriately accept or override AI triage output, not raw model accuracy [39]. This is consistent with the only close structural analogue identified in the current literature, a peer-support system proposing autonomous handling for low-stress input, moderator review for moderate-stress, and direct human intervention for high-stress cases [40] — independent convergence on the same three-tier logic from a different research group.

5.3 False Positives, False Negatives, and Why Calibration Is the Answer, Not Accuracy

Fifty years of suicide-risk-prediction research have not moved meaningfully past chance-level accuracy [18], and models with strong aggregate accuracy have been shown to perform near zero at predicting an actual future event, a direct consequence of suicide’s low base rate — not a fixable data or model-size problem [19]. This is why Sentinel’s detection layer is specified around conformal-calibrated prediction sets [24] rather than point accuracy: it is designed to communicate uncertainty honestly rather than imply a precision the underlying statistics cannot support. A January 2026 scoping review found crisis-escalation protocols “inconsistently specified” across the AI-mental-health field even where privacy and bias were addressed [37], underscoring that this is an open field-wide problem this paper is explicitly trying to close, not a gap unique to this draft.

5.4 Interpretive Robustness and Cultural Validity

Apparent empathy and correct interpretation are separable, and only the latter should be treated as safety-relevant. Current large language models, despite guardrails, still generate inaccurate or unsafe advice in a minority of self-harm scenarios [36], and sustained sycophantic validation has been documented to worsen delusional and manic states in vulnerable users [31]. Separately, culturally coded

expressions of distress — idioms that do not map cleanly onto Western diagnostic constructs [32, 33] — risk systematic misinterpretation by a system tuned on largely Western/English training data, a concern directly relevant given Tele MANAS’s own 20-language operating scope [14, 34, 35]. Section 4.3’s Stage 2 red-team battery exists specifically to test this axis before deployment.

5.5 Informed Consent, Transparency, and User Control

Three commitments follow directly from the above rather than existing as boilerplate: (1) onboarding must disclose, in plain language, exactly which signals are collected and what the escalation thresholds do, not only that “AI monitoring” is active; (2) the user retains standing control to view, adjust, or disable any escalation tier — including the Tier 3 pre-authorized contact — at any time, not only at initial setup; (3) the system discloses its non-human status at every active-conversation turn, consistent with recommendations against anthropomorphism-driven overtrust in this literature [29].

5.6 Risk of Misuse

A pre-authorized emergency contact is structurally identical to features already documented as vectors for coercive control by intimate partners [8]. The proposed mitigation is a cooling-off period before a contact can be changed, and a flag if the same contact is repeatedly added and removed.

5.7 Limitations

Three limitations are structural, not incidental. First, suicide-risk prediction is bounded by base-rate statistics fifty years of research have not moved past chance-level on [18, 19]; the conformal-calibration layer is a response to that constraint, not a claim to have solved it. Second, this paper presents a research architecture and staged safety case, not a validated clinical system — Stage 3 clinical calibration under IRB oversight is identified future work, and no result here should be read as clinical-grade until it is complete. Third, the on-device generative component inherits known large-language-model failure modes that retrieval-augmented grounding and EmoGuard reduce but do not eliminate [36]; Stage 2’s interpretive-robustness battery exists to surface these failures pre-deployment rather than assume them away. A narrower fourth limitation: passive-sensing reliability depends on Android’s background-execution policy, which can throttle precisely the low-engagement pattern the system aims to detect (Section 4.1).

6. Conclusion and Future Work

This paper set out to accomplish three things: present an architecture for fully local, multimodal mental-health anomaly detection; examine its computational feasibility under genuine mobile hardware constraints; and propose an ethical framework addressing false-positive harm, cultural and interpretive validity, and the risk of protocol misuse. Section 3 delivered the architecture — a two-phase detection pipeline wrapped in conformal calibration, a generation layer grounded in clinical instruments rather than open retrieval, and hardware-encrypted local storage with an explicit, single, user-authorized exception for crisis escalation. Section 4 delivered the feasibility analysis, including the platform-level constraints — standby-bucket throttling, KV-cache scaling — that a purely architectural description would have left unstated. Section 5 delivered the ethical framework, anchored in the closest available precedent for this exact problem, the Canada Protocol [48], and extended with a clinician-review pathway, an interpretive-robustness battery answering directly to documented sycophancy and cross-cultural misreading risks, and a staged validation plan that treats calibrated uncertainty as the honest alternative to a precision claim the underlying statistics cannot support.

That last point is the paper’s central and most load-bearing claim, restated once more because it governs how every other claim here should be read: fifty years of suicide-risk research have not moved meaningfully past chance-level prediction [18], and this is a property of the base rate, not a deficiency correctable by a better model [19]. Sentinel’s contribution is not a claim to have solved that problem. It is a proposed architecture and staged safety case for building a detection system that is honest about that constraint by design — through conformal calibration rather than point accuracy, through a human-anchored escalation protocol rather than autonomous action, and through a validation plan that names Stage 3 clinical calibration as necessary, unfinished, future work rather than eliding it.

Future work, concretely:

1. Stage 3 clinical validation is the single most important open item: an IRB-approved longitudinal pilot under licensed clinician oversight, without which no claim in this paper should be read as clinical-grade.
2. Primary-source verification of the citation on near-zero suicide-event predictive accuracy [19], used throughout this paper but not yet independently confirmed against its original publication.
3. Extension of the interpretive-robustness battery to additional Indian languages beyond current benchmark coverage, given Tele MANAS’s 20-language operating scope and the demonstrated risk of ethnocentric misreading in this literature [34, 35].
4. Direct regulatory engagement with FDA’s Digital Health Center of Excellence and India’s CDSCO, replacing the secondary regulatory analysis in Section 4.4 with primary guidance specific to a real submission.

5. Multimodal extension of the evaluation framework to voice tone and non-verbal cues, a direction the Thera-Turing Test’s own authors identify as unaddressed in their current framework [12].

6. Formal adoption of CONSORT-AI/SPIRIT-AI reporting [50, 51] once Stage 1–2 empirical results exist to report, rather than only architectural specification.

7. Empirical benchmarking of the full on-device model stack — latency, memory, battery, thermal behavior — on representative mid-range Android hardware, closing the gap this paper has been explicit about throughout between proposed design and demonstrated performance.

If these steps are completed and the staged gates are met rather than assumed, the architecture proposed here offers a plausible path toward passive mental-health monitoring that does not require choosing between clinical usefulness and zero-knowledge privacy. It does not, on its own, offer more than that — and the paper is written to make that boundary legible rather than to obscure it.

References

- [1] Torous J, Kiang MV, Lorme J, Onnela J-P. New tools for new research in psychiatry: a scalable and customizable platform to empower data driven smartphone research (introduces "digital phenotyping" and the Beiwe platform). *JMIR Ment Health*. 2016;3(2):e16.
- [2] World Health Organization. World Mental Health Today / Mental Health Atlas 2024, and Suicide Worldwide in 2021: Global Health Estimates. WHO, 2025.
- [3] Saeb S, Zhang M, Karr CJ, Schueller SM, Corden ME, Kording KP, Mohr DC. Mobile phone sensor correlates of depressive symptom severity in daily-life behavior: an exploratory study. *J Med Internet Res*. 2015;17(7):e175.
- [4] Wang R, Chen F, Chen Z, Li T, Harari G, Tignor S, Zhou X, Ben-Zeev D, Campbell AT. StudentLife: assessing mental health, academic performance and behavioral trends of college students using smartphones. *Proc. ACM UbiComp 2014*, pp. 3-14.
- [5] Vaidyam A, Halamka J, Torous J. Enabling research and clinical use of patient-generated health data (the mindLAMP platform): digital phenotyping study. *JMIR Mhealth Uhealth*. 2022;10(1):e30557.
- [6] McMahan HB, Moore E, Ramage D, Hampson S, y Arcas BA. Communication-efficient learning of deep networks from decentralized data. *Proc. AISTATS 2017*, pp. 1273-1282.
- [7] Triaging casual from critical: leveraging machine learning to detect self-harm and suicide risks for youth on social media. *JMIR Ment Health*, 2026. (Full author byline to be confirmed against primary source before submission.)
- [8] Freed D, Palmer J, Minchala D, Levy K, Ristenpart T, Dell N. "A Stalker's Paradise": how intimate partner abusers exploit technology. *Proc. ACM CHI 2018*, pp. 1-13.
- [9] Android Open Source Project / Android Developers. Hardware-backed Keystore documentation (Trusted Execution Environment, StrongBox). [source.android.com](https://source.android.com/docs/security/keystore/hw-backed-keystore); [developer.android.com](https://developer.android.com/docs/security/keystore/hw-backed-keystore).
- [10] Kehinde O, Adeyeye E, Ayanwale MA, et al. Can AI express empathy? Computational insights from mental health dialogues. *J Technol Behav Sci*. 2026.
- [11] EmoAgent: assessing and safeguarding human-AI interaction for mental health safety (introduces the EmoEval and EmoGuard components). *arXiv:2504.09689*, 2025 / *EMNLP 2025*.
- [12] Bunge EL, Desage C. A framework for evaluating mental health artificial intelligence-based conversational agents (introduces the Thera-Turing Test). *J Technol Behav Sci*. 2025;10:731-739.
- [13] McBain RK, Cantor JH, Breslau J, et al. AI chatbot use and disclosure for mental health among US adolescents and young adults. *JAMA Pediatr*. 2026.
- [14] Ministry of Health and Family Welfare, Government of India / NIMHANS. Tele MANAS (Tele Mental Health Assistance and Networking Across States). Launched October 10, 2022. telemanas.mohfw.gov.in.
- [15] Sun Z, Yu H, Song X, Liu R, Yang Y, Zhou D. MobileBERT: a compact task-agnostic BERT for resource-limited devices. *Proc. ACL 2020*.
- [16] Warner S, et al. Smarter, Better, Faster, Longer: A Modern Bidirectional Encoder for Fast, Memory Efficient, and Long Context Finetuning and Inference (ModernBERT). *arXiv:2412.13663*, 2024.
- [17] Fitzpatrick KK, Darcy A, Vierhile M. Delivering cognitive behavior therapy to young adults with symptoms of depression and anxiety using a fully automated conversational agent (Woebot): a randomized controlled trial. *JMIR Ment Health*. 2017;4(2):e19.

- [18] Franklin JC, Ribeiro JD, Fox KR, Bentley KH, Kleiman EM, Huang X, Musacchio KM, Jaroszewski AC, Chang BP, Nock MK. Risk factors for suicidal thoughts and behaviors: a meta-analysis of 50 years of research. *Psychol Bull.* 2017;143(2):187-232.
- [19] Belsher BE, et al. Prediction models for suicide attempts and deaths: a systematic review and simulation. (Primary citation to be independently verified before submission; referenced here via secondary literature describing near-chance predictive accuracy for actual suicide events despite high aggregate classification accuracy.)
- [20] Cao B, Zulueta J, Leow A, et al. Predicting mood disturbance severity with mobile phone keystroke metadata: a BiAffect digital phenotyping study. *J Med Internet Res.* 2018;20(7):e241.
- [21] Vesel C, et al. Effects of mood and aging on keystroke dynamics metadata in a large open-science sample: a BiAffect iOS study. *J Am Med Inform Assoc.* 2020;27(7):1007-1018.
- [22] Lewis P, Perez E, Piktus A, et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Proc. NeurIPS* 2020.
- [23] Xu, Yan, Dai, Wu. MEGA-RAG: a retrieval-augmented generation framework with multi-evidence guided answer refinement for mitigating hallucinations of LLMs in public health. *Front Public Health.* 2025;13:1635381.
- [24] Conformal learning for quantifying uncertainty in psychiatric condition predictions (NIH All of Us Fitbit-derived data, 13,835 participants). Preprint, 2025.
- [25] Google. Sleep API (Activity Recognition APIs). Google Play Services documentation.
- [26] Google DeepMind. Gemma 4 model family (E2B/E4B) and on-device deployment guidance, including the June 2026 quantization-aware training release. ai.google.dev/gemma.
- [27] Google. UsageStatsManager. Android Developers reference documentation.
- [28] Stade EC, Stirman SW, Ungar LH, Boland CL, Schwartz HA, Yaden DB, Sedoc J, DeRubeis R, Willer R, Eichstaedt JC. Large language models could change the future of behavioral healthcare: a proposal for responsible development and evaluation. *NPJ Ment Health Res.* 2024;3(1):12.
- [29] De Freitas J, Uğuralp AK, Oğuz-Uğuralp Z, Puntoni S. Chatbots and mental health: insights into the safety of generative AI. *J Consum Psychol.* 2024;34(3):481-491.
- [30] Lawrence HR, Schneider RA, Rubin SB, Matarić MJ, McDuff DJ, Bell MJ. The opportunities and risks of large language models in mental health. *JMIR Ment Health.* 2024;11(1):e59479.
- [31] Technological folie à deux: feedback loops between AI chatbots and mental illness. *arXiv:2507.19218*, 2025.
- [32] Kleinman A. Depression, somatization and the "new cross-cultural psychiatry." *Soc Sci Med.* 1977;11:3-9.
- [33] Kohrt BA, Hruschka DJ. Nepali concepts of psychological trauma: the role of idioms of distress. *Cult Med Psychiatry.* 2010;34:322-352.
- [34] PsychiatryBench: a multi-task benchmark for LLMs in psychiatry. *arXiv:2509.09711*, 2025.
- [35] Cross-lingual mental health ontologies for Indian languages: bridging patient expression and clinical understanding through explainable AI. *arXiv:2510.05387*, 2025.
- [36] Kolding S, et al. LLM guardrail performance in self-harm scenarios, 2024. (Primary citation to be independently verified before submission.)
- [37] Scoping review of reviews: artificial intelligence in mental health care. *Front Psychiatry.* 2026.
- [38] EthicsMH: a pilot benchmark for ethical reasoning in mental health AI. *arXiv:2509.11648*, 2025.

- [39] SimClinician: a multimodal simulation testbed for reliable psychologist-AI collaboration in mental health diagnosis. *arXiv:2512.08953*, 2025.
- [40] Investigating AI in peer support via multi-module system-driven embodied conversational agents. *arXiv:2511.22269*, 2025.
- [41] FDA-authorized software as a medical device in mental health: a perspective on evidence, device lineage, and regulatory challenges. *NPJ Ment Health Res.* 2025.
- [42] Linardon J, Fuller-Tyszkiewicz M. Attrition and adherence in smartphone-delivered interventions for mental health problems: a systematic and meta-analytic review. *J Consult Clin Psychol.* 2020;88(1):1-13.
- [43] U.S. Food and Drug Administration. Artificial Intelligence/Machine Learning (AI/ML)-Based Software as a Medical Device (SaMD) Action Plan. January 2021.
- [44] Rosenthal JT, Beecy A, Sabuncu MR. Rethinking clinical trials for medical AI with dynamic deployments of adaptive systems. *NPJ Digit Med.* 2025;8:252.
- [45] Livingston L, Featherstone-Uwague A, Barry A, Barretto K, Morey T, Herrmannova D, Avula V. Reproducible generative AI evaluation for healthcare: a clinician-in-the-loop approach. *medRxiv* 2025.03.04.25323131.
- [46] Corfmat M, Martineau JT, Régis C. High-reward, high-risk technologies? An ethical and legal account of AI development in healthcare. *BMC Med Ethics.* 2025;26:4.
- [47] Saeidnia HR, Hashemi Fotami SG, Lund B, Ghiasi N. Ethical considerations in artificial intelligence interventions for mental health and well-being: ensuring responsible implementation and impact. *Soc Sci.* 2024;13:381.
- [48] Mörch C-M, Gupta A, Mishara BL. Canada Protocol: an ethical checklist for the use of artificial intelligence in suicide prevention and mental health. *Artif Intell Med.* 2020;108:101934.
- [49] Dehbozorgi R, Zangeneh S, Khooshab E, Hafezi Nia D, Hanif HR, Samian P, Yousefi M, Haj Hashemi F, Vakili M, Jamalimoghadam N, Lohrasebi F. The application of artificial intelligence in the field of mental health: a systematic review. *BMC Psychiatry.* 2025;25:132.
- [50] Liu X, Cruz Rivera S, Moher D, Calvert MJ, Denniston AK. Reporting guidelines for clinical trial reports for interventions involving artificial intelligence: the CONSORT-AI extension. *Lancet Digit Health.* 2020;2:e537-e548.
- [51] Cruz Rivera S, Liu X, Chan A-W, Denniston AK, Calvert MJ. Guidelines for clinical trial protocols for interventions involving artificial intelligence: the SPIRIT-AI extension. *Lancet Digit Health.* 2020;2:e549-e560.
- [52] Roy K, Gaur M, Soltani M, Rawte V, Kalyan A, Sheth A. ProKnow: process knowledge for safety constrained and explainable question generation for mental health diagnostic assistance. *Front Big Data.* 2023;5:1056728.
- [53] Kellogg KC, Sadeh-Sharvit S. Pragmatic AI-augmentation in mental healthcare: key technologies, potential benefits, and real-world challenges and solutions for frontline clinicians. *Front Psychiatry.* 2022;13:990370.
- [54] Amugongo LM, Mascheroni P, Brooks S, Doering S, Seidel J. Retrieval augmented generation for large language models in healthcare: a systematic review. *PLOS Digit Health.* 2025;4(6):e0000877.